

Microsoft Recommenders

Using Jupyter notebooks, Python and the <https://www.github.com/Microsoft/Recommenders> repo, I'll demonstrate how you can make your own Recommendation system using AzureML. I'll note the benefits of Azure ML and its ability to use all of the usual Python tools and libraries that are popular today. In addition, various models used for recommendation systems will be compared using predictions developed on the MovieLens dataset.

<https://github.com/Microsoft/Recommenders>
<https://notebooks.azure.com/>



Learn to make your own Recommendation System with AzureML

Beth Zeranski

Principal Experimentation Engineer

<http://github.com/bethz/2019GlobalAzureBootcamp>

<https://github.com/Microsoft/Recommenders>

<https://twitter.com/bethzeranski>

Linkedin: [aka.ms/bethz](https://www.linkedin.com/company/aka.ms/bethz)

April 27, 2019

Agenda

- Goal
- Why Recommendation Systems?
- What is a Recommendation System?
 - Powerpoint's *AWESOME* Recommendation System
- <http://github.com/Microsoft/Recommenders>
- Solution Architecture
- Data – public movielens data
- Walk through Repo
- Azure Machine Learning, AzureML
- Workflow for SAR, Smart Adaptive Recommendation solution
 - <https://github.com/Microsoft/Product-Recommendations/blob/master/doc/sar.md>
- Code Walk through

Goal

Provide sufficient context and guidance so you can use the Recommenders github repo to make your own Recommendation System

Highlight notebooks demonstrate multiple algorithms, automated parameter tuning and even algorithm selection to assist.



<https://github.com/Microsoft/Recommenders>

Huge Revenue Impact!

Why Recommendation Systems?

A collage of four images. The top-left image shows a woman with blonde hair in a braid, wearing a purple lab coat and a headset with an orange cord, looking at a device in a laboratory setting. The top-right image shows a man with grey hair and glasses, wearing a white t-shirt, working in a field of green plants. The bottom-left image shows a man in a red shirt pointing at a large monitor displaying a 3D brain scan. The bottom-right image shows a man with a beard and sunglasses, smiling while holding a white smartphone. The text is overlaid on the top-left image.

“35% of what consumers purchase on Amazon and 75% of what they watch on Netflix come from recommendations algorithms”

McKinsey & Co

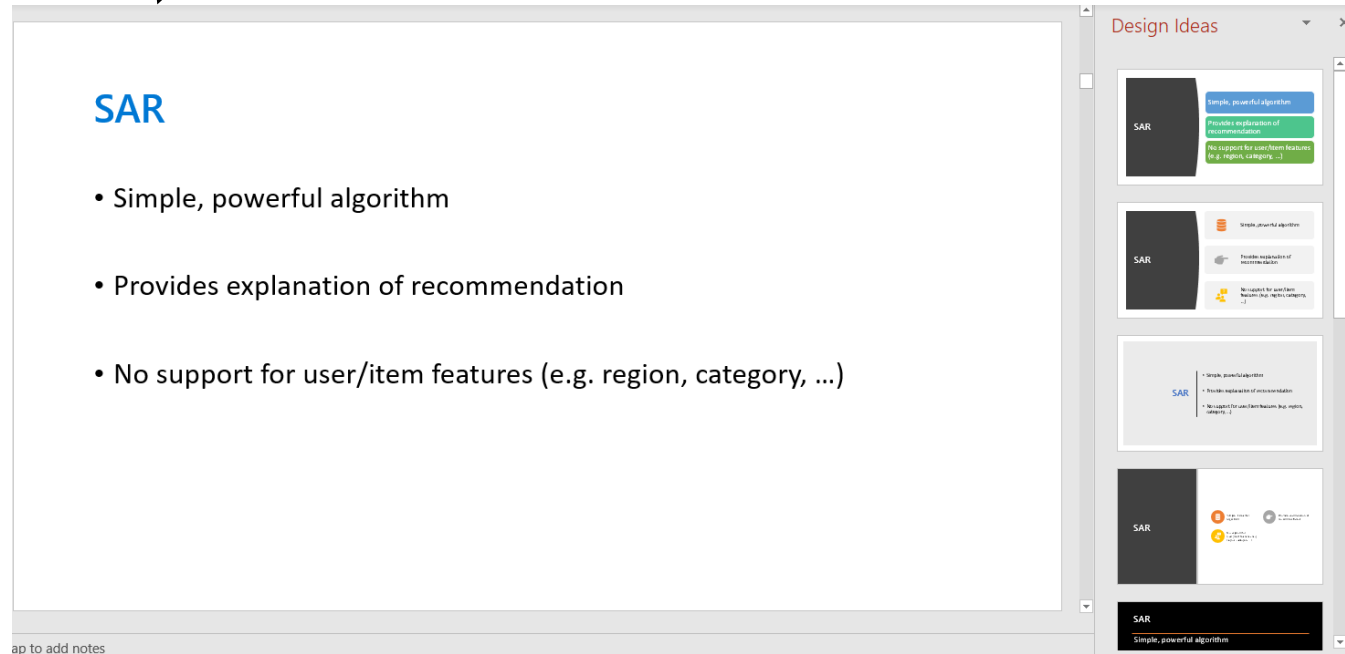
What is a Recommendation System?

...is a subclass of information filtering systems that seek to predict the 'rating' or 'preference' that a user would give to an item

https://en.wikipedia.org/wiki/Recommender_system

Powerpoint has an
AWESOME
Recommendation
System

My simple
slide



The screenshot shows a PowerPoint slide titled "SAR" with the following bullet points:

- Simple, powerful algorithm
- Provides explanation of recommendation
- No support for user/item features (e.g. region, category, ...)

To the right of the slide is the "Design Ideas" pane, which displays several alternative slide layouts for the same content. Each layout in the pane shows a different arrangement of the text and bullet points, some with additional visual elements like icons or background colors. The pane is titled "Design Ideas" and has a close button (X) in the top right corner.

Design
Ideas

- Powerpoint's "Design Ideas" Recommendation system created a number of improved layouts which are more compelling for my slides
- Fast and Easy

Key Themes of our *Recommenders repo*

- Reduce time and effort required for Proof of Concept
- Lower barrier of entry for implementation
- Easily scale out on Azure
- Widely usable and open-source
- Provide advanced algorithms

The screenshot shows the GitHub interface for the Microsoft/Recommenders repository. The repository is private and has 18 watchers, 7 unstars, and 4 forks. The current view is the 'Code' tab, showing a file named 'sar_python_cpu_movielens.ipynb' in the '00_quick_start' directory. The file was committed by miguelgferro and has 4 contributors. The file size is 23.5 KB and it contains 732 lines of code. The notebook content is displayed, showing the title 'SAR on MovieLens (Python, CPU)' and the following text:

SAR is a fast scalable adaptive algorithm for personalized recommendations based on user transaction history. It produces easily explainable / interpretable recommendations and handles "cold item" and "semi-cold user" scenarios. SAR is a kind of neighborhood based algorithm (as discussed in [Recommender Systems by Aggarwal](#)) which is intended for ranking top items for each user.

SAR recommends items that are most **similar** to the ones that the user already has an existing **affinity** for. Two items are **similar** if the users who have interacted with one item are also likely to have interacted with another. A user has an **affinity** to an item if they have interacted with it in the past.

Advantages of SAR:

- High accuracy for an easy to train and deploy algorithm
- Fast training, only requiring simple counting to construct matrices used at prediction time.
- Fast scoring, only involving multiplication of the similarity matrix with an affinity vector

Notes to use SAR properly:

- Since it does not use item or user features, it can be at a disadvantage against algorithms that do.
- It's memory-hungry, requiring the creation of an $m \times m$ sparse square matrix (where m is the number of items). This can also be a problem for many matrix factorization algorithms.
- SAR favors an implicit rating scenario and it does not predict ratings.

This notebook provides an example of how to utilize and evaluate SAR in Python on a CPU.

0 Global Settings and Imports

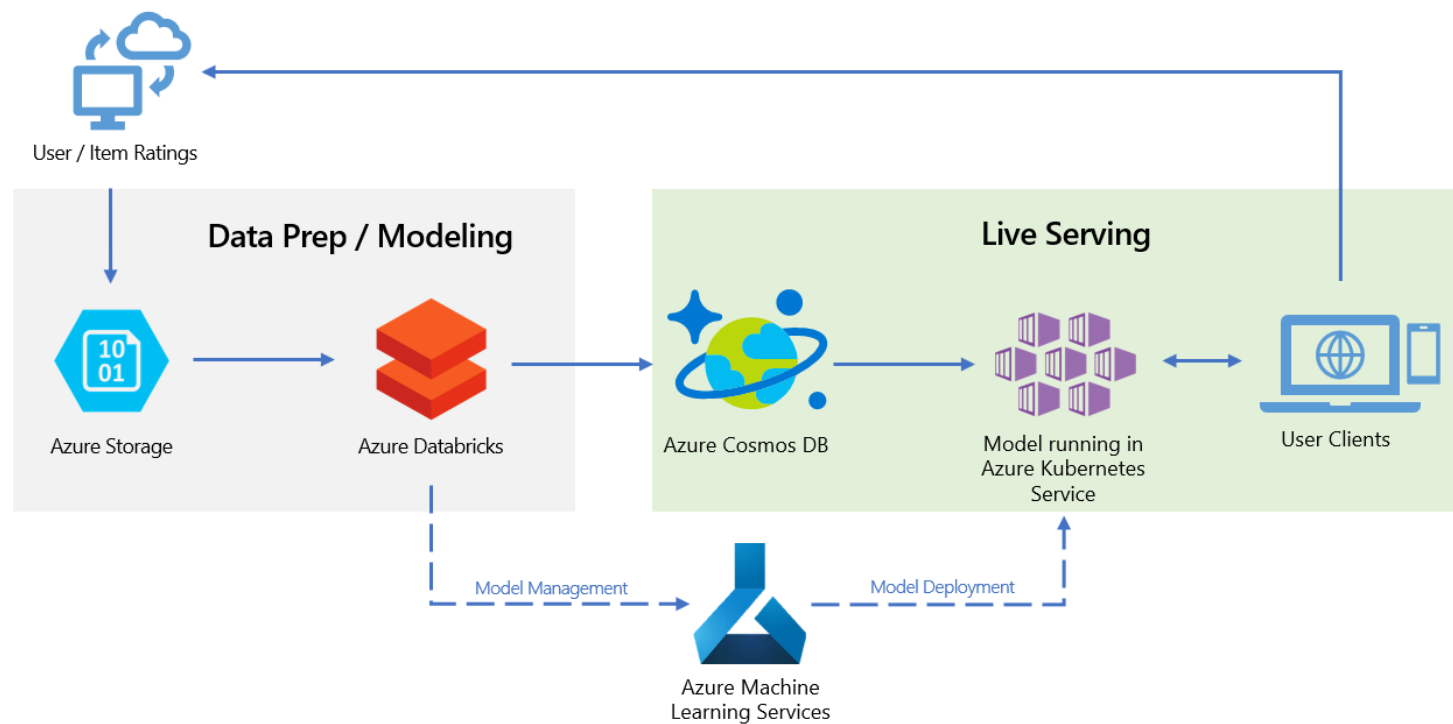
```
In [2]: # set the environment path to find Recommenders
import sys
sys.path.append("../..")
```

<https://github.com/Microsoft/Recommenders>

Microsoft/Recommenders

- Microsoft/Recommenders
 - Collaborative development efforts of Microsoft Cloud & AI data scientists, Microsoft Research researchers, academia researchers, etc.
 - Github url: <https://github.com/Microsoft/Recommenders>
 - Contents
 - Utilities: modular functions for model creation, data manipulation, evaluation, etc.
 - Algorithms: SVD, SAR, ALS, NCF, Wide&Deep, xDeepFM, DKN, etc.
 - Notebooks: HOW-TO examples for end to end recommender building.
 - Highlights
 - 2800+ stars on GitHub
 - Featured in YC Hacker News, O'Reilly Data Newsletter, GitHub weekly trending list, etc.
 - Any contribution to the repo will be highly appreciated!
 - Create issue/PR directly in the GitHub repo
 - Send email to RecoDevTeam@service.microsoft.com for any collaboration

Recommenders Architecture



UserId	MovieId	Rating	Timestamp
196	242	3.0	881250949
186	302	3.0	891717742
22	377	1.0	878887116
244	51	2.0	880606923
166	346	1.0	886397596
298	474	4.0	884182806
115	265	2.0	881171488

- User ID
- Item ID
- Time (Optional)
- Event Weight (Optional)
 - Rating
 - Number of item views
 - Number of purchases

Data format


Repository General Components

Reco Utilities




Dataset helpers
Splitters
Ranking Metrics
Rating Metrics
Diversity Metrics

Notebooks



Spark ALS
SAR: Spark and Python
implementation
Deep Learning algos from MSR
(DeepRec)
Open source frameworks: Surprise

Tests



Unit tests
Integration tests
Nightly builds
Benchmarks updated every
night

Repository Walk-thru

Microsoft / **Recommenders**

Unwatch

120

★ Unstar

3,009

Fork

381

<> Code

! Issues 74

Pull requests 10

Projects 2

Wiki

Insights

Best Practices on Recommendation Systems

machine-learning

recommender

ranking

deep-learning

python

jupyter-notebook

recommendation-algorithm

rating

operationalization

kubernetes

azure

microsoft

2,110 commits

27 branches

3 releases

30 contributors

MIT

Branch: master

New pull request

Create new file

Upload files

Find File

Clone or download

miguelgferro Merge pull request #720 from Microsoft/staging ...

Latest commit 1b682fc 8 days ago

.github	Update text	2 months ago
benchmarks	addressed @anargyri comments 🚀	11 days ago
notebooks	Merge pull request #687 from Microsoft/andreas/nni	8 days ago
reco_utils	Merge pull request #687 from Microsoft/andreas/nni	8 days ago
scripts	Merge branch 'staging' into andreas/nni	10 days ago
tests	Merge pull request #687 from Microsoft/andreas/nni	8 days ago

<https://github.com/Microsoft/Recommenders>

Why Azure ML Service?

Training in the cloud

Highly scalable and flexible model training

Rapid Iteration

Model Traceability

Automated Data, Output, Notebook and related Asset Retention

Azure ML's automated machine learning finds algorithms based on your data

Pipelines for repeatability and sharing

Workspace

- The workspace is the top-level resource for Azure Machine Learning service. It provides a centralized place to work with all the artifacts you create when you use Azure Machine Learning service.
- The workspace keeps a list of compute targets that you can use to train your model. It also keeps a history of the training runs, including logs, metrics, output, and a snapshot of your scripts. You use this information to determine which training run produces the best model.
- You register models with the workspace. You use a registered model and scoring scripts to create an image. You can then deploy the image to Azure Container Instances, Azure Kubernetes Service, or to a field-programmable gate array (FPGA) as a REST-based HTTP endpoint. You can also deploy the image to an Azure IoT Edge device as a module.
- You can create multiple workspaces, and each workspace can be shared by multiple people.
- Workspace is accessible in portal.azure.com or by Command Line Interface

For more info:

- <https://docs.microsoft.com/en-us/azure/machine-learning/service/concept-azure-machine-learning-architecture>

Experiment

- An experiment is a grouping of many runs from a specified script. It always belongs to a workspace. When you submit a run, you provide an experiment name. Information for the run is stored under that experiment. If you submit a run and specify an experiment name that doesn't exist, a new experiment with that newly specified name is automatically created.
- For an example of using an experiment, see [Quickstart: Get started with Azure Machine Learning service](#).

Model

- At its simplest, a model is a piece of code that takes an input and produces output. Creating a machine learning model involves selecting an algorithm, providing it with data, and tuning hyperparameters. Training is an iterative process that produces a trained model, which encapsulates what the model learned during the training process.
- A model is produced by a run in Azure Machine Learning. You can also use a model that's trained outside of Azure Machine Learning. You can register a model in an Azure Machine Learning service workspace.
- Azure Machine Learning service is framework agnostic. When you create a model, you can use any popular machine learning framework, such as Scikit-learn, XGBoost, PyTorch, TensorFlow, and Chainer.
- For an example of training a model, see [Tutorial: Train an image classification model with Azure Machine Learning service.](#)


Model registry

- The model registry keeps track of all the models in your Azure Machine Learning service workspace.
- Models are identified by name and version. Each time you register a model with the same name as an existing one, the registry assumes that it's a new version. The version is incremented, and the new model is registered under the same name.
- When you register the model, you can provide additional metadata tags and then use the tags when you search for models.
- You can't delete models that are being used by an image.
- For an example of registering a model, see [Train an image classification model with Azure Machine Learning](#).

Run Configuration



An environment defined for each compute target so that you can run your training scripts on different compute targets *without* modification.



For example, to use a local target configuration:

```
from azureml.core import ScriptRunConfig
import os
script_folder = os.getcwd()
src = ScriptRunConfig(source_directory = script_folder,
script = 'train.py', run_config = run_local)
run = exp.submit(src)
run.wait_for_completion(show_output = True)
```

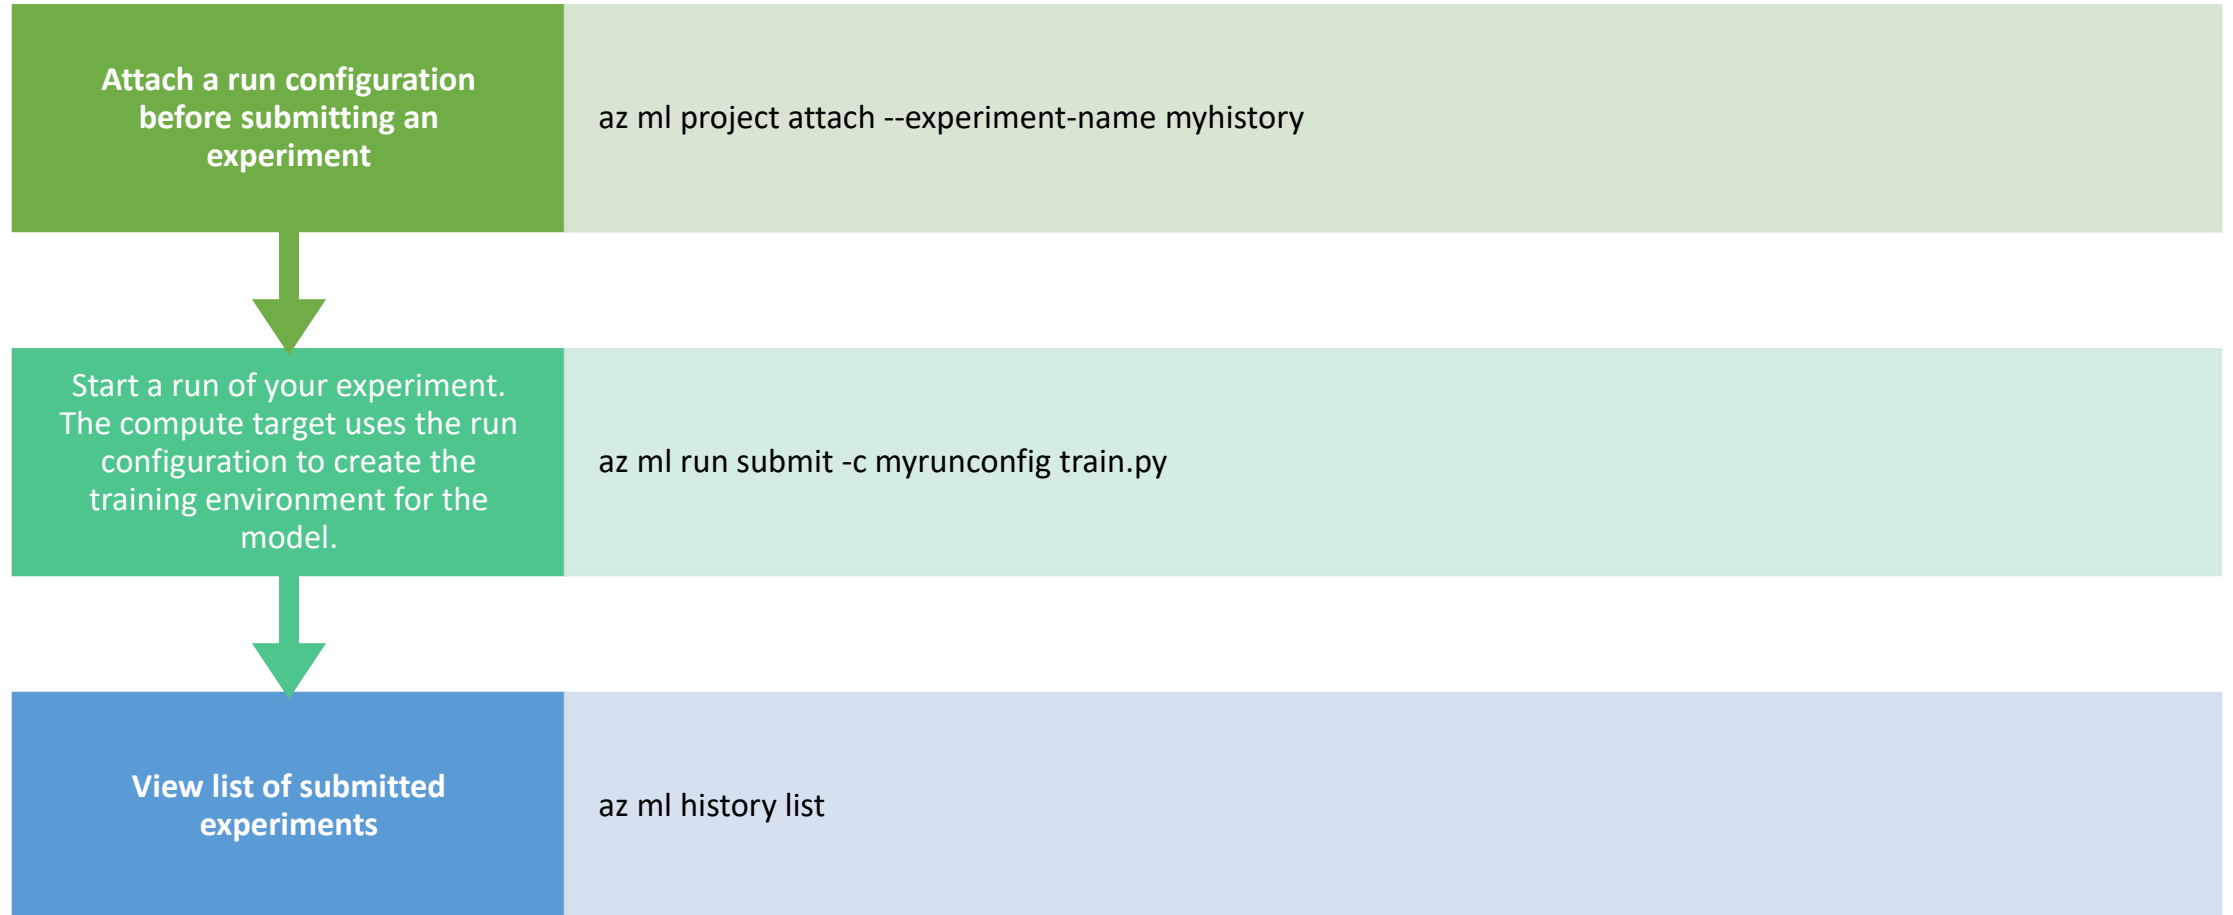


Switch the same experiment to run in a different compute target by using a different run configuration, such as the AzureML Compute target:

```
from azureml.core import ScriptRunConfig

src = ScriptRunConfig(source_directory = script_folder,
script = 'train.py', run_config = run_amlcompute)
run = exp.submit(src)
run.wait_for_completion(show_output = True)
```

Command Line Interface



During local
training or
cloud training,
have you ever
lost a run?

- Days or weeks go by and models+info go missing
- Azure ML service enables you to log artifacts from local experimentation or cloud training painlessly to the cloud. That includes inputs, notebooks, outputs and more!
- That means the work is easily repeatable (*and findable*) by you or anyone with access

Other info

- Exporting and deleting data
- <https://docs.microsoft.com/en-us/azure/machine-learning/service/how-to-export-delete-data>
- Start, monitor and cancel training runs
- <https://docs.microsoft.com/en-us/azure/machine-learning/service/how-to-manage-runs>
- Log Metrics
- <https://docs.microsoft.com/en-us/azure/machine-learning/service/how-to-track-experiments>
- Estimator – convenient object that wraps run configuration information to specify how a script is executed for model training
- <https://github.com/Azure/MachineLearningNotebooks/blob/master/how-to-use-azureml/training-with-deep-learning/how-to-use-estimator/how-to-use-estimator.ipynb>

Workflow

The machine learning workflow generally follows this sequence:

1. Develop machine learning training scripts in **Python**.
2. Create and configure a **compute target**.
3. Attach the compute target to your workspace.
4. Configure the compute target so it contains the Python environment and package dependencies required.
5. **Submit the scripts** to the configured compute target to run in that environment. During training, the scripts can read from or write to **datastore**. The records of execution are saved as **runs** in the **workspace** and grouped under **experiments**.
6. **Query the experiment** for logged metrics from the current and past runs. If the metrics don't indicate a desired outcome, loop back to step 1 and iterate on your scripts. All artifacts are stored together for revisiting prior work and outcomes.
7. After a satisfactory run is identified, register the persisted model in the **model registry**.
8. Develop a scoring script.
9. **Create an image** and register it in the **image registry**.
10. **Deploy the image** as a **web service** in Azure.

Start Here

- <https://github.com/Azure/MachineLearningNotebooks/blob/master/configuration.ipynb>

Demo: SAR Quickstart

https://github.com/Microsoft/Recommenders/blob/master/notebooks/00_quick_start/sar_movielens_with_azureml.ipynb

https://github.com/Microsoft/Recommenders/blob/master/notebooks/run_notebook_on_azureml.ipynb

https://github.com/Microsoft/Recommenders/blob/master/notebooks/00_quick_start/sar_movielens.ipynb

References

Great detail on Algorithm and deeper dive on creating a system

www.github.com/yueguoguo/pakdd2019tutorial

Train an image classification model with AzureML

<https://github.com/Azure/MachineLearningNotebooks/blob/master/tutorials/img-classification-part1-training.ipynb>

Deploy an image classification model in Azure Container Instance

<https://github.com/Azure/MachineLearningNotebooks/blob/master/tutorials/img-classification-part2-deploy.ipynb>

Textbook:

Recommender Systems: The Textbook by Charu C. Aggarwal

<https://www.amazon.com/Recommender-Systems-Textbook-Charu-Aggarwal/dp/3319296574>

Blog:

<https://azure.microsoft.com/en-us/blog/experimentation-using-azure-machine-learning/>

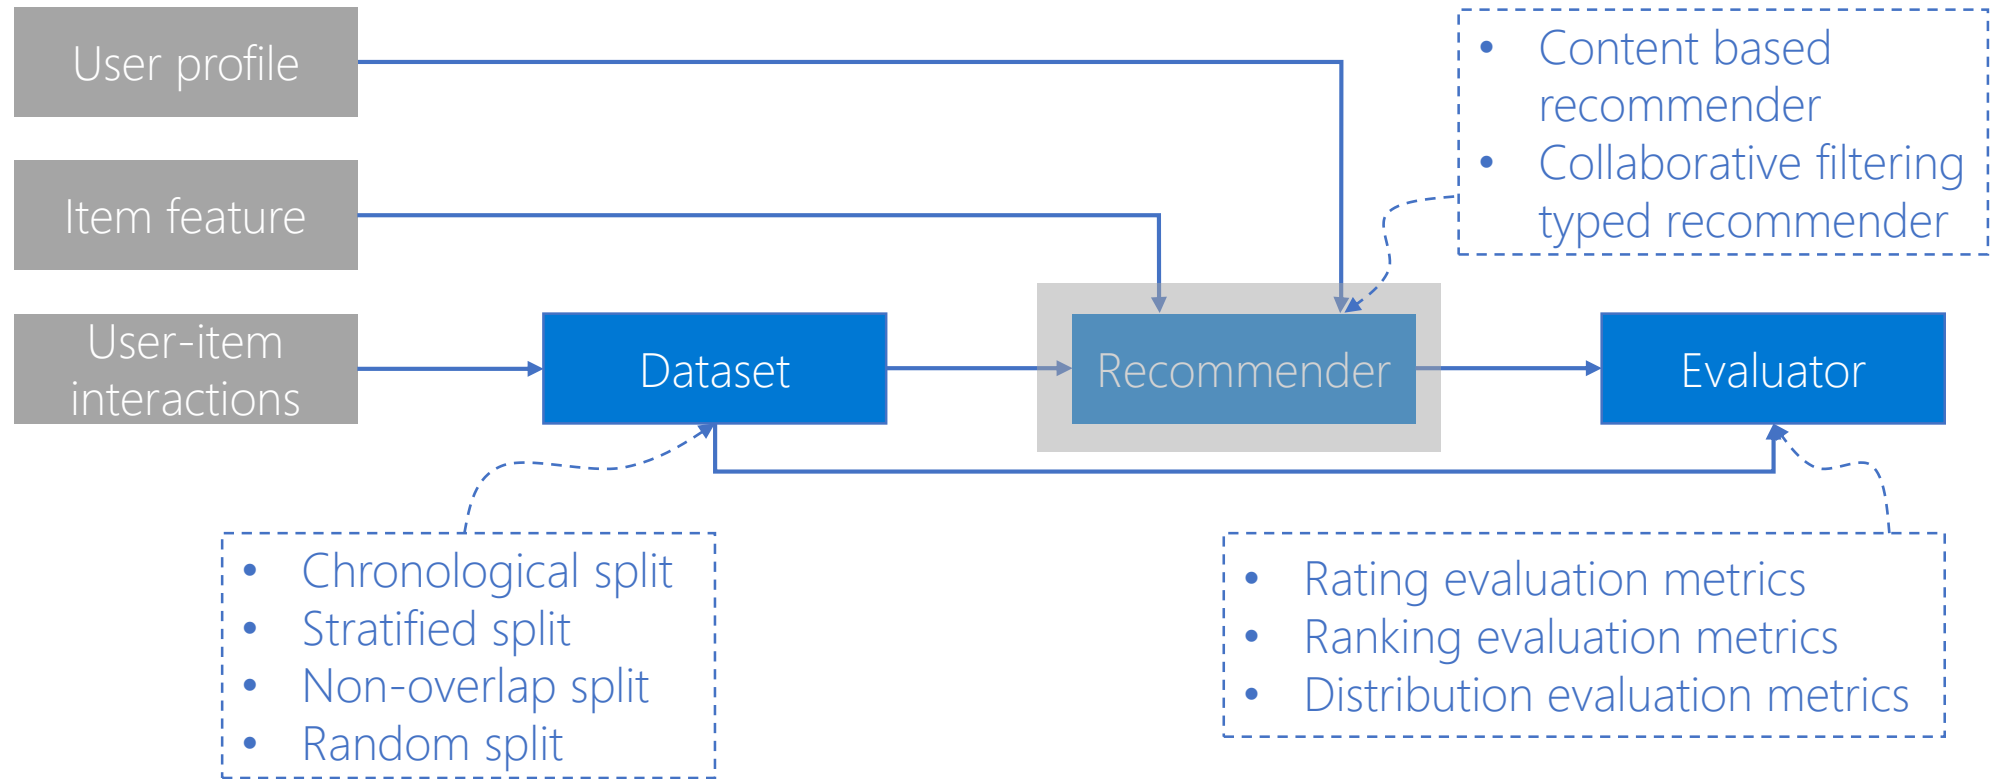


Questions?

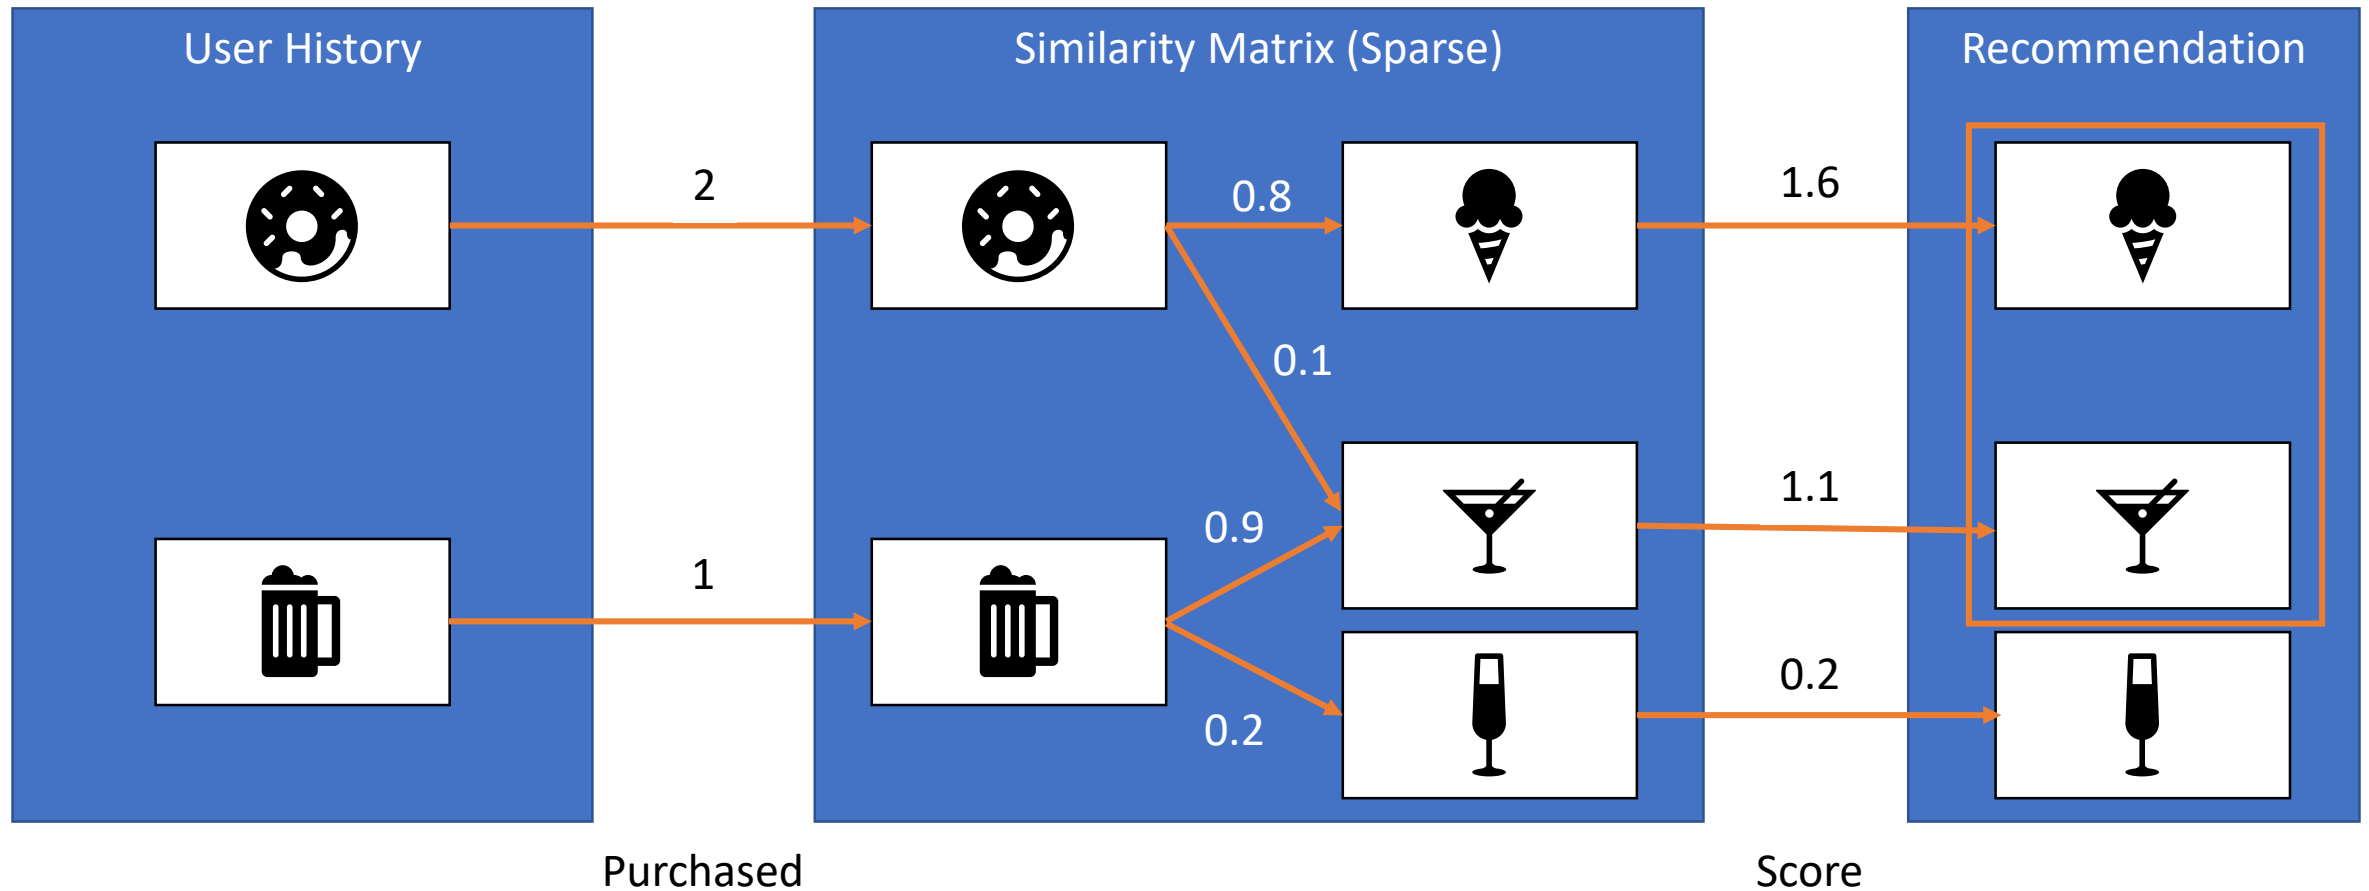
Microsoft Recommenders

<https://github.com/Microsoft/Recommenders>

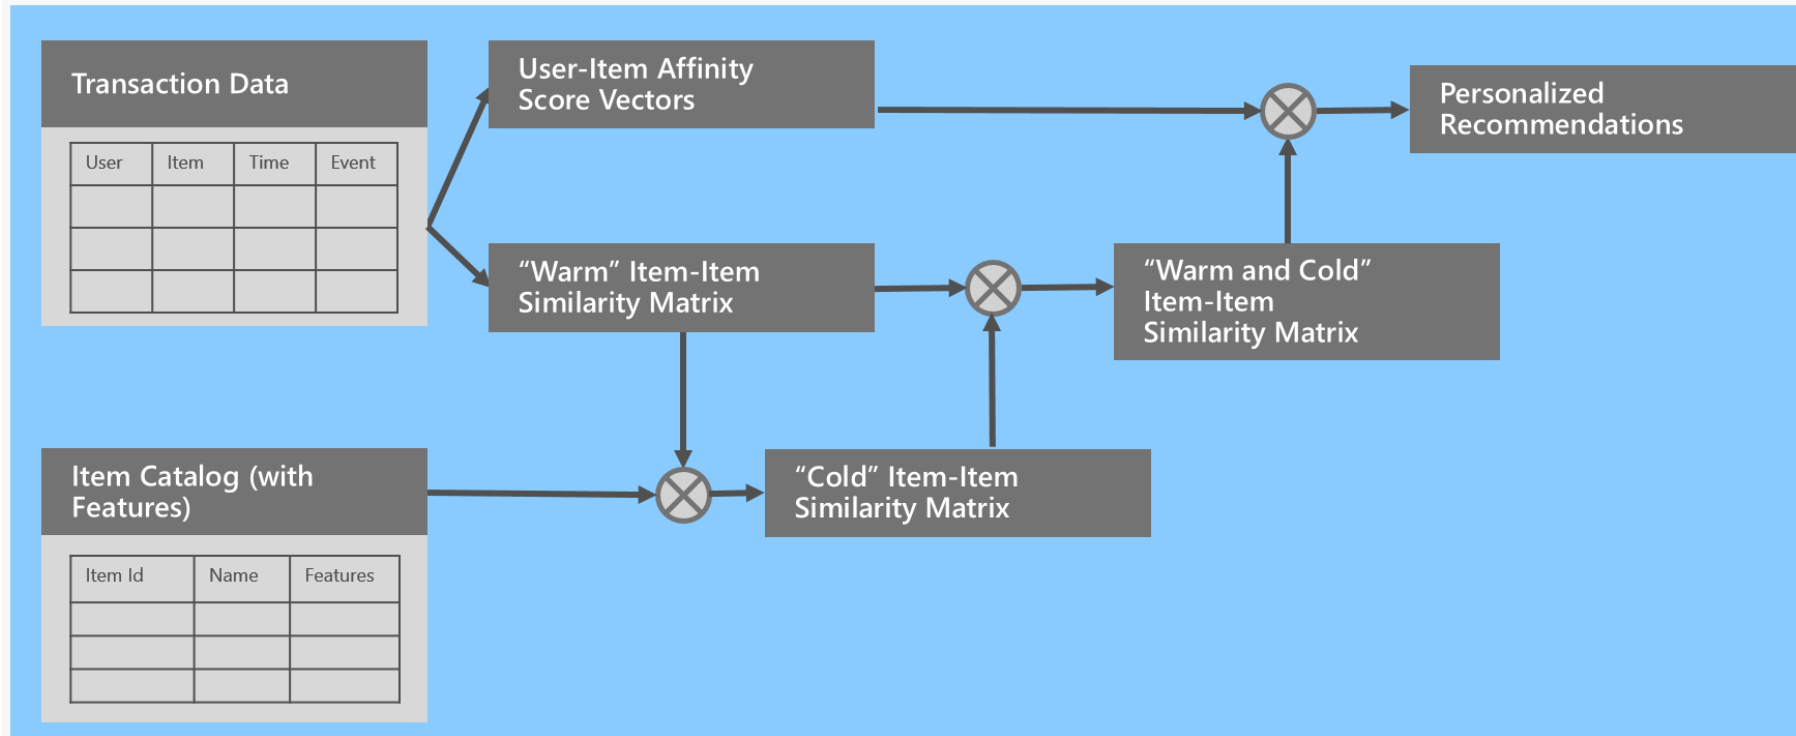
Example of an end-to-end pipeline



SAR



SAR overview



Data Splitting

- Stratified splitting
 - e.g. movie ratings
 - Avoid cold-users
 - Train users = test users
- Chronological splitting
 - e.g. fashion items, news, ...
 - User taste changes over time
 - Obey time-dependency

SAR

- Simple, powerful algorithm
- Provides explanation of recommendation
- No support for user/item features (e.g. region, category, ...)